

**Sri Sivasubramaniya Nadar College of Engineering, Chennai**  
(An autonomous Institution affiliated to Anna University)

Degree & Branch	B.E. Computer Science & Engineering	Semester	V
Subject Code & Name	ICS1512 & Machine Learning Algorithms Laboratory		
Academic year	2025-2026 (Odd)	Batch:2023-2028	<b>Due date:</b>

## Experiment 2 : Loan Amount Prediction using Linear Regression

### 1. Aim

To perform end-to-end preprocessing, visualization, and model training on a real-world loan prediction dataset using Python libraries such as **pandas**, **numpy**, **matplotlib**, **seaborn**, and **scikit-learn**. The task includes cleaning the data, encoding and normalizing features, applying regression, and evaluating model performance using cross-validation.

### 2. Libraries Used

- **NumPy**: For performing numerical computations and handling arrays efficiently.
  - `np.mean()`, `np.std()`, `np.where()`, `np.log1p()`
- **Pandas**: For loading, cleaning, and transforming tabular data.
  - `pd.read_csv()`, `df.drop()`, `df.isna()`, `df.replace()`, `df.describe()`
- **Matplotlib**: For visualizing data through static plots.
  - `plt.plot()`, `plt.scatter()`, `plt.barh()`, `plt.title()`, `plt.show()`
- **Seaborn**: For advanced statistical visualizations.
  - `sns.histplot()`, `sns.boxplot()`, `sns.heatmap()`, `sns.scatterplot()`
- **scikit-learn (sklearn)**: For preprocessing, regression modeling, and evaluation.
  - `LinearRegression`, `train_test_split()`, `KFold()`, `cross_validate()`, `OneHotEncoder`, `ColumnTransformer`, `LabelEncoder`, `StandardScaler`, `SimpleImputer`
  - `mean_absolute_error()`, `mean_squared_error()`, `r2_score()`
- **Google Colab**: Used as the execution environment for the experiment.
  - `files.upload()`, `!kaggle datasets download`, `!unzip`

### 3. Objective

#### 4.1 Loading and Cleaning the Dataset

- Unnecessary columns like **Name** and **Customer ID** were dropped.
- Missing values were handled using:

- Median imputation for numerical features.
- Mode imputation for categorical features.
- Target column (**Loan Sanction Amount (USD)**) rows with missing values were removed for training.

### 3.2 Feature Engineering

- Binary categorical columns (**Expense Type 1**, **Expense Type 2**) were mapped to 0/1.
- One-hot encoding was applied to categorical features using **ColumnTransformer** and **OneHotEncoder** from **scikit-learn**.

### 3.3 Exploratory Data Analysis (EDA)

- Histograms, boxplots, and scatter plots were used to understand the distribution and relationships of features.
- Log-transformations were applied to skewed features like **Income**.
- Correlation heatmaps were used to identify features with strong correlation to the target.

### 3.4 Outlier Handling and Normalization

- Outliers were capped using the IQR (Interquartile Range) method.
- Z-score normalization was applied to numerical columns including the target variable.

### 3.5 Model Building and Evaluation

- A **LinearRegression** model was trained on the normalized dataset.
- Performance metrics calculated:
  - Mean Absolute Error (MAE)
  - Root Mean Squared Error (RMSE)
  - Coefficient of Determination ( $R^2$ )
- A residual plot was used to check the assumptions of linear regression.
- Feature importance was visualized using the learned coefficients.

### 3.6 Cross-Validation

- K-Fold Cross-Validation ( $K = 5$ ) was applied to estimate generalization error.
- Fold-wise metrics were recorded and tabulated for analysis.

## 4. Mathematical Description of Linear Regression

Linear Regression is a supervised learning algorithm used for predicting a continuous dependent variable ( $y$ ) based on one or more independent variables ( $x_1, x_2, \dots, x_n$ ).

## 4.1 Hypothesis Function

The hypothesis function for multiple linear regression is defined as:

$$h_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n = \theta^T \mathbf{x}$$

Where:

- $\mathbf{x} = [1, x_1, x_2, \dots, x_n]^T$  is the feature vector (with 1 for intercept)
- $\theta = [\theta_0, \theta_1, \dots, \theta_n]^T$  are the model parameters
- $h_{\theta}(\mathbf{x})$  is the predicted output

## 4.2 Cost Function

The cost function used to minimize the error in prediction is the **\*\*Mean Squared Error (MSE)\*\***:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

Where:

- $m$  is the number of training examples
- $h_{\theta}(x^{(i)})$  is the predicted value for the  $i^{th}$  training example
- $y^{(i)}$  is the actual output

## 4.3 Gradient Descent Update Rule

To minimize the cost function, we update the parameters using gradient descent:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

The partial derivative of the cost function with respect to  $\theta_j$  is:

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)}$$

Hence, the update rule becomes:

$$\theta_j := \theta_j - \alpha \cdot \frac{1}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)}$$

Where  $\alpha$  is the learning rate.

## 5. Code with plots

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, KFold, cross_validate
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Upload kaggle.json
from google.colab import files
files.upload() # Select your kaggle.json file

# Rename and move the file properly
!mkdir -p ~/.kaggle
!mv "kaggle (3).json" ~/.kaggle/kaggle.json
!chmod 600 ~/.kaggle/kaggle.json

# Install Kaggle CLI (if not already installed)
!pip install -q kaggle

# Download the dataset
!kaggle datasets download -d phileinsophos/predict-loan-amount-data

# Unzip the downloaded file
!unzip -o predict-loan-amount-data.zip
```

<IPython.core.display.HTML object>

Saving kaggle (3).json to kaggle (3).json

Dataset URL: <https://www.kaggle.com/datasets/phileinsophos/predict-loan-amount-data>

License(s): CC0-1.0

Downloading predict-loan-amount-data.zip to /content

0% 0.00/2.44M [00:00<?, ?B/s]

100% 2.44M/2.44M [00:00<00:00, 544MB/s]

Archive: predict-loan-amount-data.zip

inflating: test.csv

inflating: train.csv

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, KFold, cross_validate
from sklearn.preprocessing import LabelEncoder, StandardScaler
```

```

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# My Dataset
df = pd.read_csv('train.csv')
print(df.columns)

```

```

Index(['Customer ID', 'Name', 'Gender', 'Age', 'Income (USD)',
      'Income Stability', 'Profession', 'Type of Employment', 'Location',
      'Loan Amount Request (USD)', 'Current Loan Expenses (USD)',
      'Expense Type 1', 'Expense Type 2', 'Dependents', 'Credit Score',
      'No. of Defaults', 'Has Active Credit Card', 'Property ID',
      'Property Age', 'Property Type', 'Property Location', 'Co-Applicant',
      'Property Price', 'Loan Sanction Amount (USD)'],
      dtype='object')

```

## 2. Data Preprocessing

```

[3]: import pandas as pd
from sklearn.impute import SimpleImputer

# 2. Data Preprocessing
# ----- Remove unnecessary columns ----- #

df.drop(columns=['Name', 'Customer ID'], inplace=True)

print(df.columns)
# ----- Handling missing values ----- #
# replace unknown values with np.nan

df = df.replace(r'(?i)^\s*unknown\s*$', np.nan, regex=True)

print(df.isna().sum())

# Define numerical columns (excluding target)
num_cols = ['Income (USD)', 'Credit Score', 'Current Loan Expenses (USD)',
            'Dependents', 'Property Age']

# Define categorical columns
cat_cols = ['Gender', 'Income Stability', 'Type of Employment',
            'Has Active Credit Card', 'Property Location']

# -----
# Impute numerical with MEDIAN
# -----
num_imputer = SimpleImputer(strategy='median')
df[num_cols] = num_imputer.fit_transform(df[num_cols])

```

```

# -----
# Impute categorical with MODE
# -----
cat_imputer = SimpleImputer(strategy='most_frequent')
df[cat_cols] = cat_imputer.fit_transform(df[cat_cols])

# -----
# Handle target variable (Loan Sanction Amount)
# -----

# Leave missing target values as-is - you'll train only on non-missing rows
print(" Missing values in target before training:", df['Loan Sanction Amount_
↳(USD)'].isnull().sum())

#For training later:
df = df[df['Loan Sanction Amount (USD)'].notnull()]
print(df.columns)

```

```

Index(['Gender', 'Age', 'Income (USD)', 'Income Stability', 'Profession',
      'Type of Employment', 'Location', 'Loan Amount Request (USD)',
      'Current Loan Expenses (USD)', 'Expense Type 1', 'Expense Type 2',
      'Dependents', 'Credit Score', 'No. of Defaults',
      'Has Active Credit Card', 'Property ID', 'Property Age',
      'Property Type', 'Property Location', 'Co-Applicant', 'Property Price',
      'Loan Sanction Amount (USD)'],
      dtype='object')

```

Gender	53
Age	0
Income (USD)	4576
Income Stability	1683
Profession	0
Type of Employment	7270
Location	0
Loan Amount Request (USD)	0
Current Loan Expenses (USD)	172
Expense Type 1	0
Expense Type 2	0
Dependents	2493
Credit Score	1703
No. of Defaults	0
Has Active Credit Card	1566
Property ID	0
Property Age	4850
Property Type	0
Property Location	356
Co-Applicant	0
Property Price	0

```

Loan Sanction Amount (USD)      340
dtype: int64
Missing values in target before training: 340
Index(['Gender', 'Age', 'Income (USD)', 'Income Stability', 'Profession',
      'Type of Employment', 'Location', 'Loan Amount Request (USD)',
      'Current Loan Expenses (USD)', 'Expense Type 1', 'Expense Type 2',
      'Dependents', 'Credit Score', 'No. of Defaults',
      'Has Active Credit Card', 'Property ID', 'Property Age',
      'Property Type', 'Property Location', 'Co-Applicant', 'Property Price',
      'Loan Sanction Amount (USD)'],
      dtype='object')

```

```

[4]: import pandas as pd
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
# Convert 'y' and 'n' to 1 and 0
df['Expense Type 1'] = df['Expense Type 1'].map({'Y': 1, 'N': 0})
df['Expense Type 2'] = df['Expense Type 2'].map({'Y': 1, 'N': 0})

# Define categorical columns
categorical_cols = [
    'Gender',
    'Income Stability',
    'Profession',
    'Type of Employment',
    'Location',
    'Has Active Credit Card',
    'Property Location'
]

# Separate features and target if needed (assume target not encoded)
X = df.copy()

# Define ColumnTransformer with OneHotEncoder (drop='first' to avoid
↳ multicollinearity)
column_transformer = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(drop='first', sparse_output=False),
↳ categorical_cols)
    ],
    remainder='passthrough' # keep other columns
)

# Fit and transform the data
X_encoded_array = column_transformer.fit_transform(X)

```

```

# Get new column names from the encoder
encoded_cat_cols = column_transformer.named_transformers_['cat'].
    ↳get_feature_names_out(categorical_cols)

# Create DataFrame with encoded data
df_encoded = pd.DataFrame(X_encoded_array, columns=encoded_cat_cols.tolist() +
    ↳[col for col in X.columns if col not in categorical_cols])

# Optional: Convert all columns to appropriate types
df_encoded = df_encoded.convert_dtypes()

# Reset index if needed
df_encoded.index = df.index

# Output
print(df_encoded.head())
print(df_encoded.columns)
df = df_encoded
print(df.columns)
print("Combined shape:", df.shape)
print("Encoded shape:", df.shape)
print("Memory used:", df.memory_usage(deep=True).sum() / (1024**2), "MB")

```

	Gender_M	Income	Stability_Low	Profession_Commercial associate \
0	0		1	0
1	1		1	0
2	0		0	0
3	0		0	0
4	0		1	0

	Profession_Maternity leave	Profession_Pensioner	Profession_State servant \
0	0	0	0
1	0	0	0
2	0	1	0
3	0	1	0
4	0	0	0

	Profession_Student	Profession_Unemployed	Profession_Working \
0	0	0	1
1	0	0	1
2	0	0	0
3	0	0	0
4	0	0	1

	Type of Employment_Cleaning staff ...	Expense Type 2	Dependents \
0	0 ...	0	3
1	0 ...	1	1



2	0	...	1	1
3	0	...	1	2
4	0	...	1	2

	Credit Score	No. of Defaults	Property ID	Property Age	Property Type \
0	809.44	0	746	1933.05	4
1	780.4	0	608	4952.91	2
2	833.15	0	546	988.19	2
3	832.7	1	890	2223.25	2
4	745.55	1	715	2614.77	4

	Co-Applicant	Property Price	Loan Sanction Amount (USD)
0	1	119933.46	54607.18
1	1	54791.0	37469.98
2	0	72440.58	36474.43
3	1	121441.51	56040.54
4	1	208567.91	74008.28

[5 rows x 47 columns]

```
Index(['Gender_M', 'Income Stability_Low', 'Profession_Commercial associate',
      'Profession_Maternity leave', 'Profession_Pensioner',
      'Profession_State servant', 'Profession_Student',
      'Profession_Unemployed', 'Profession_Working',
      'Type of Employment_Cleaning staff', 'Type of Employment_Cooking staff',
      'Type of Employment_Core staff', 'Type of Employment_Drivers',
      'Type of Employment_HR staff',
      'Type of Employment_High skill tech staff',
      'Type of Employment_IT staff', 'Type of Employment_Laborers',
      'Type of Employment_Low-skill Laborers', 'Type of Employment_Managers',
      'Type of Employment_Medicine staff',
      'Type of Employment_Private service staff',
      'Type of Employment_Realty agents', 'Type of Employment_Sales staff',
      'Type of Employment_Secretaries', 'Type of Employment_Security staff',
      'Type of Employment_Waiters/barmen staff', 'Location_Semi-Urban',
      'Location_Urban', 'Has Active Credit Card_Inactive',
      'Has Active Credit Card_Unpossessed', 'Property Location_Semi-Urban',
      'Property Location_Urban', 'Age', 'Income (USD)',
      'Loan Amount Request (USD)', 'Current Loan Expenses (USD)',
      'Expense Type 1', 'Expense Type 2', 'Dependents', 'Credit Score',
      'No. of Defaults', 'Property ID', 'Property Age', 'Property Type',
      'Co-Applicant', 'Property Price', 'Loan Sanction Amount (USD)'],
      dtype='object')
Index(['Gender_M', 'Income Stability_Low', 'Profession_Commercial associate',
      'Profession_Maternity leave', 'Profession_Pensioner',
      'Profession_State servant', 'Profession_Student',
      'Profession_Unemployed', 'Profession_Working',
      'Type of Employment_Cleaning staff', 'Type of Employment_Cooking staff',
      'Type of Employment_Core staff', 'Type of Employment_Drivers',
```

```

'Type of Employment_HR staff',
'Type of Employment_High skill tech staff',
'Type of Employment_IT staff', 'Type of Employment_Laborers',
'Type of Employment_Low-skill Laborers', 'Type of Employment_Managers',
'Type of Employment_Medicine staff',
'Type of Employment_Private service staff',
'Type of Employment_Realty agents', 'Type of Employment_Sales staff',
'Type of Employment_Secretaries', 'Type of Employment_Security staff',
'Type of Employment_Waiters/barmen staff', 'Location_Semi-Urban',
'Location_Urban', 'Has Active Credit Card_Inactive',
'Has Active Credit Card_Unpossessed', 'Property Location_Semi-Urban',
'Property Location_Urban', 'Age', 'Income (USD)',
'Loan Amount Request (USD)', 'Current Loan Expenses (USD)',
'Expense Type 1', 'Expense Type 2', 'Dependents', 'Credit Score',
'No. of Defaults', 'Property ID', 'Property Age', 'Property Type',
'Co-Applicant', 'Property Price', 'Loan Sanction Amount (USD)'],
dtype='object')
Combined shape: (29660, 47)
Encoded shape: (29660, 47)
Memory used: 12.19125747680664 MB

```

```

[5]: import matplotlib.pyplot as plt
import seaborn as sns

# List of features
features = ['Age', 'Income (USD)', 'Loan Amount Request (USD)',
           'Current Loan Expenses (USD)', 'Credit Score',
           'Loan Sanction Amount (USD)']

# Histograms
for col in features:
    plt.figure(figsize=(6, 4))
    sns.histplot(df[col], kde=True, bins=30, color='skyblue')
    plt.title(f'Histogram of {col}')
    plt.xlabel(col)
    plt.ylabel('Frequency')
    plt.tight_layout()
    plt.show()
    plt.close()

# Boxplots
for col in features:
    plt.figure(figsize=(6, 4))
    sns.boxplot(x=df[col], color='salmon')
    plt.title(f'Boxplot of {col}')
    plt.tight_layout()
    plt.show()


```

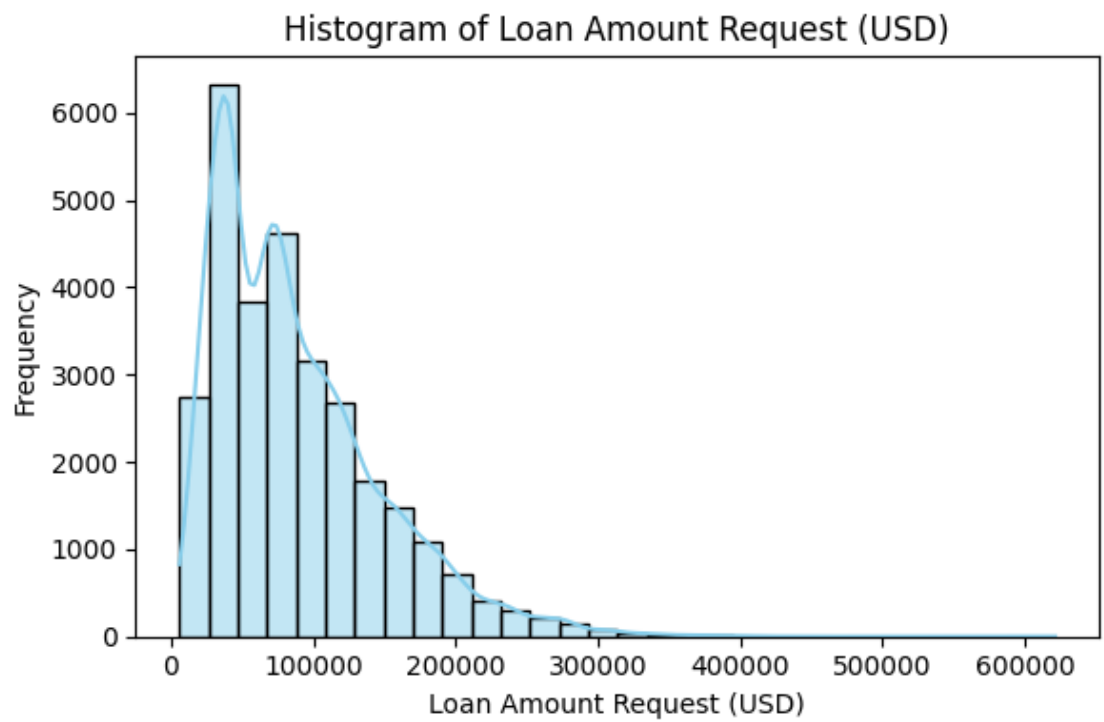

```


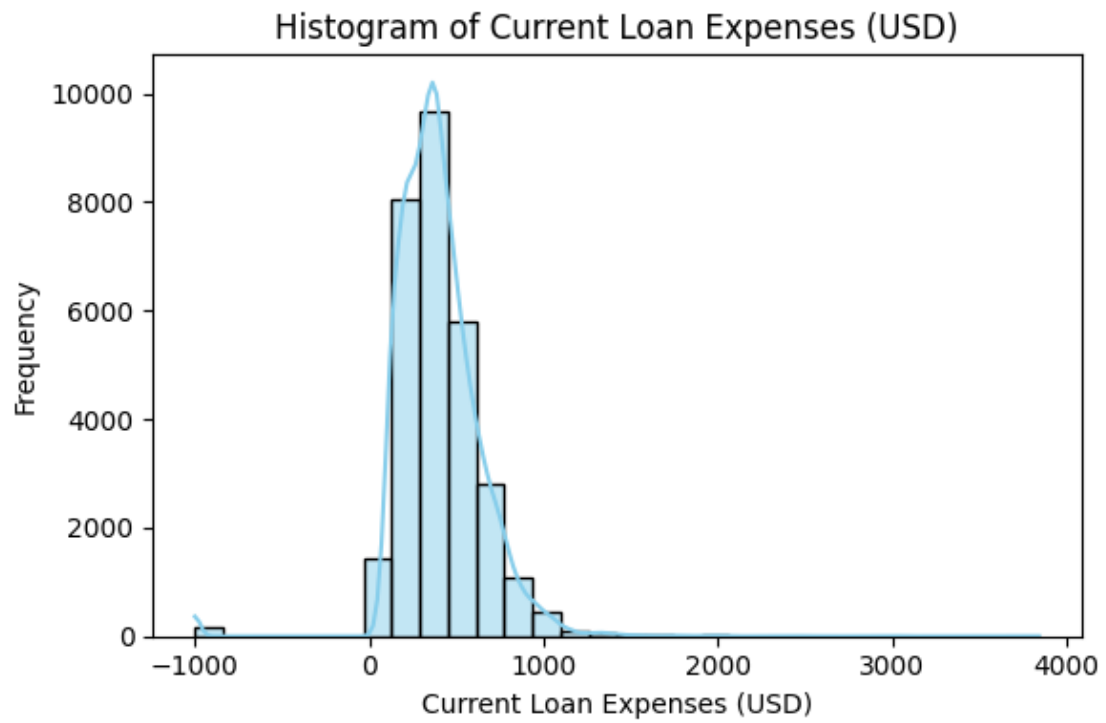
plt.close()



# Scatter plots against Loan Sanction Amount
target = 'Loan Sanction Amount (USD)'
for col in features:
    if col != target:
        plt.figure(figsize=(6, 4))
        sns.scatterplot(x=df[col], y=df[target], alpha=0.6)
        plt.title(f'{col} vs {target}')
        plt.xlabel(col)
        plt.ylabel(target)
        plt.tight_layout()
        plt.show()
        plt.close()

```









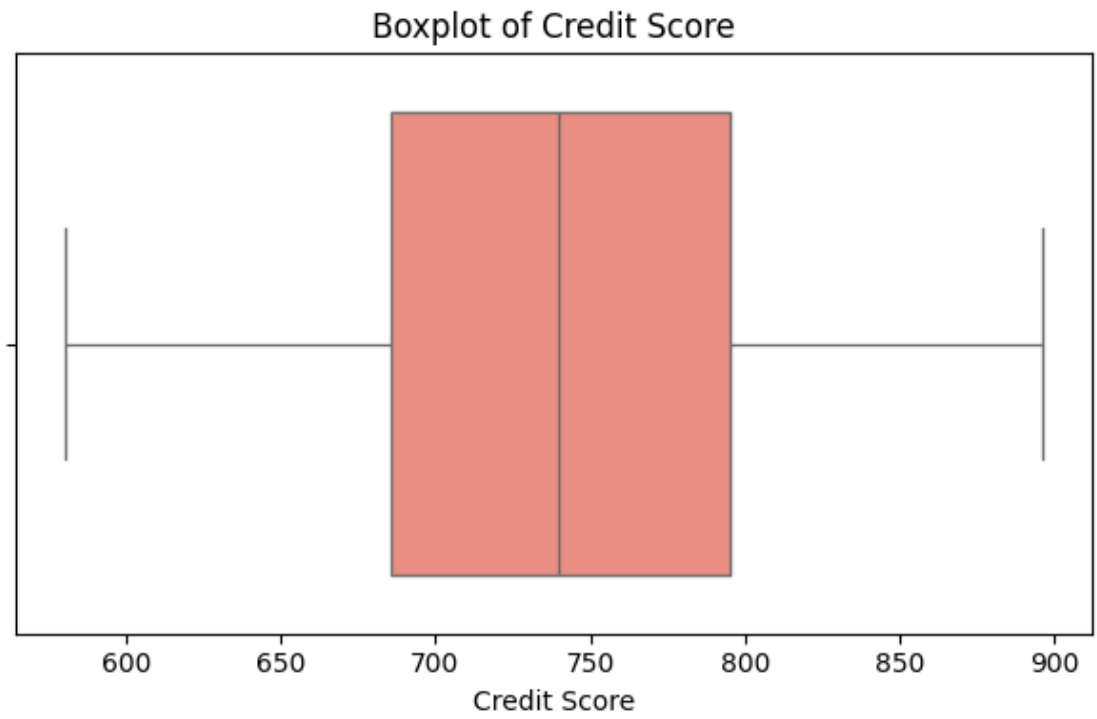
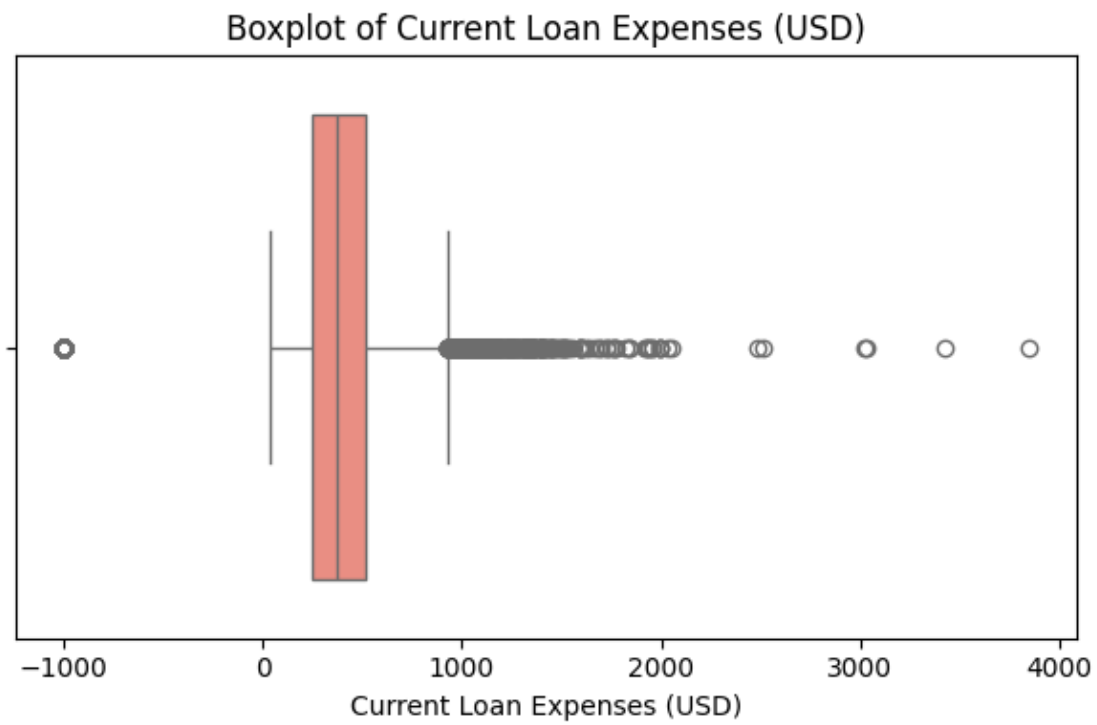


Boxplot of Income (USD)



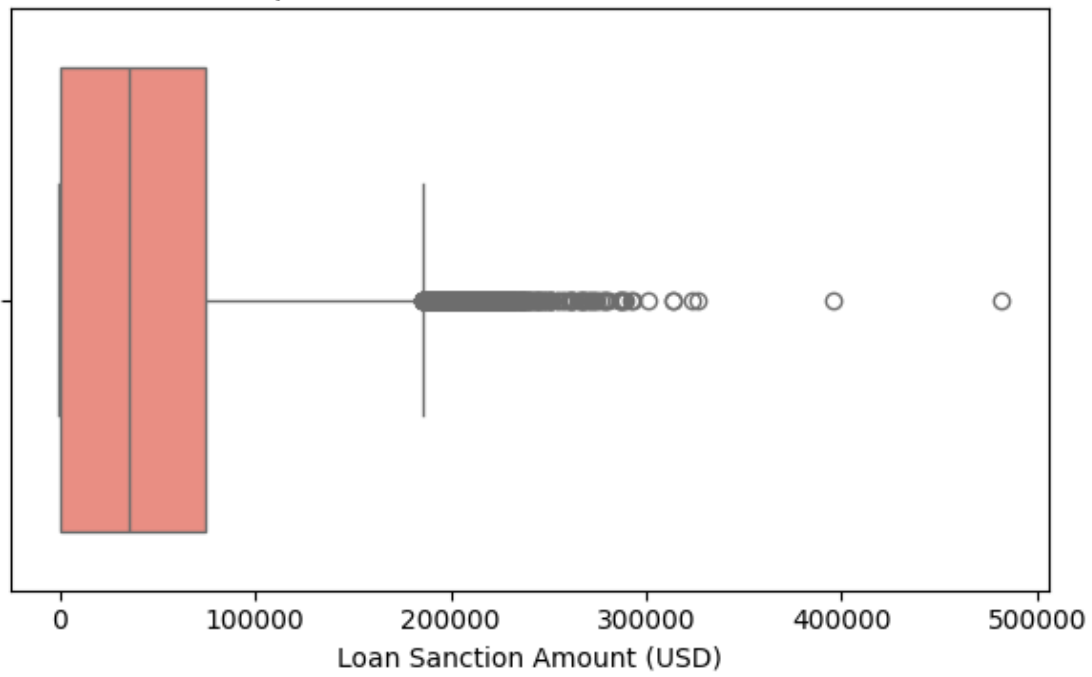
Boxplot of Loan Amount Request (USD)












Boxplot of Loan Sanction Amount (USD)



Age vs Loan Sanction Amount (USD)







```
[6]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Assuming df is already loaded and cleaned

# Feature to visualize
income_col = 'Income (USD)'
target_col = 'Loan Sanction Amount (USD)'

# Histogram (Raw)
plt.figure(figsize=(6, 4))
sns.histplot(df[income_col], kde=True, bins=50, color='skyblue')
plt.title(f'Histogram of {income_col}')
plt.xlabel(income_col)
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()
plt.close()

# Histogram (Log Transformed)
plt.figure(figsize=(6, 4))
sns.histplot(np.log1p(df[income_col]), kde=True, bins=50, color='purple')
plt.title(f'Log-Transformed Histogram of {income_col}')
plt.xlabel(f'Log({income_col} + 1)')
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()
plt.close()


# Boxplot
plt.figure(figsize=(6, 4))
sns.boxplot(x=df[income_col], color='salmon')
plt.title(f'Boxplot of {income_col}')
plt.tight_layout()
plt.show()
plt.close()



# Boxplot (Zoomed, remove extreme outliers at 99th percentile)
plt.figure(figsize=(6, 4))
filtered_income = df[df[income_col] < df[income_col].quantile(0.99)]
sns.boxplot(x=filtered_income[income_col], color='green')
plt.title(f'Boxplot of {income_col} (Zoomed)')
plt.tight_layout()
plt.show()
plt.close()


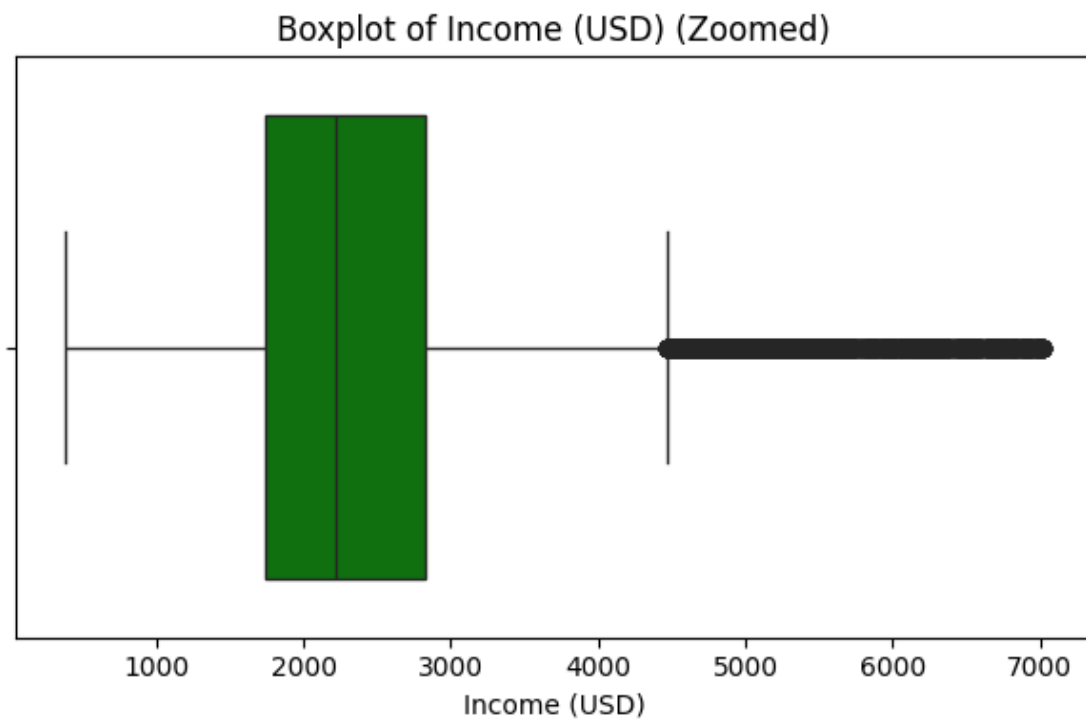
# Scatter Plot: Income vs Loan Sanction Amount
```


```
plt.figure(figsize=(6, 4))
sns.scatterplot(x=df[income_col], y=df[target_col], alpha=0.6)
plt.title(f'{income_col} vs {target_col}')
plt.xlabel(income_col)
plt.ylabel(target_col)
plt.tight_layout()
plt.show()
plt.close()

# Scatter Plot with log-transformed income
plt.figure(figsize=(6, 4))
sns.scatterplot(x=np.log1p(df[income_col]), y=df[target_col], alpha=0.6)
plt.title(f'Log({income_col} + 1) vs {target_col}')
plt.xlabel(f'Log({income_col} + 1)')
plt.ylabel(target_col)
plt.tight_layout()
plt.show()
plt.close()
```









```
[7]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Ensure only numerical columns are used for correlation
numerical_df = df.select_dtypes(include=['int64', 'float64'])

# Calculate correlations with the target column
target_col = 'Loan Sanction Amount (USD)'
correlations = {}

for col in numerical_df.columns:
    if col != target_col:
        correlations[col] = numerical_df[col].corr(numerical_df[target_col])


# Convert to DataFrame for heatmap
corr_df = pd.DataFrame.from_dict(correlations, orient='index',
    ↪columns=['Correlation with Loan Amount'])
corr_df = corr_df.sort_values(by='Correlation with Loan Amount', ascending=False)

# Plotting heatmap
plt.figure(figsize=(6, 12))
sns.heatmap(corr_df, annot=True, cmap='coolwarm', fmt=".2f", cbar=True)
```



```
plt.title(f'Correlation of Features with {target_col}')  
plt.tight_layout()  
plt.show()  
plt.close()
```

Correlation of Features with Loan Sanction Amount (USD)



Correlation with Loan Amount

## Normalization

```
[9]: import pandas as pd
import numpy as np

# Sample: Load your dataset
# df = pd.read_csv('your_dataset.csv')

# --- Function to cap outliers using IQR ---
def cap_outliers_iqr(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    df[column] = np.where(df[column] < lower_bound, lower_bound,
                          np.where(df[column] > upper_bound, upper_bound, df[column]))

# --- List of columns to cap using IQR ---
columns_to_cap = [
    'Income (USD)',
    'Loan Amount Request (USD)',
    'Current Loan Expenses (USD)',
    'Loan Sanction Amount (USD)'
]

# 1. Z-score: Income
df['Income (USD)'] = (
    df['Income (USD)'] - df['Income (USD)'].mean()
) / df['Income (USD)'].std()

# 2. Z-score: Loan Amount Request
df['Loan Amount Request (USD)'] = (
    df['Loan Amount Request (USD)'] - df['Loan Amount Request (USD)'].mean()
) / df['Loan Amount Request (USD)'].std()

# 3. Standardize: Current Loan Expenses (already Z-scored correctly)
df['Current Loan Expenses (USD)'] = (
    df['Current Loan Expenses (USD)'] - df['Current Loan Expenses (USD)'].mean()
) / df['Current Loan Expenses (USD)'].std()

# 4. Standardize: Credit Score (already Z-scored correctly)
```

```

df['Credit Score'] = (
    df['Credit Score'] - df['Credit Score'].mean()
) / df['Credit Score'].std()

# 5. Z-score: Loan Sanction Amount (Target Variable)
df['Loan Sanction Amount (USD)'] = (
    df['Loan Sanction Amount (USD)'] - df['Loan Sanction Amount (USD)'].mean()
) / df['Loan Sanction Amount (USD)'].std()

# Apply IQR capping
for col in columns_to_cap:
    cap_outliers_iqr(df, col)

# --- Show transformed data ---
print(df.head())

```

	Gender_M	Income Stability_Low	Profession_Commercial associate \
0	0	1	0
1	1	1	0
2	0	0	0
3	0	0	0
4	0	1	0

	Profession_Maternity leave	Profession_Pensioner	Profession_State servant \
0	0	0	0
1	0	0	0
2	0	1	0
3	0	1	0
4	0	0	0

	Profession_Student	Profession_Unemployed	Profession_Working \
0	0	0	1
1	0	0	1
2	0	0	0
3	0	0	0
4	0	0	1

	Type of Employment_Cleaning staff ...	Expense Type 2	Dependents \
0	0 ...	0	3
1	0 ...	1	1
2	0 ...	1	1
3	0 ...	1	2
4	0 ...	1	2

	Credit Score	No. of Defaults	Property ID	Property Age	Property Type \
0	0.992824	0	746	1933.05	4
1	0.578508	0	608	4952.91	2
2	1.331097	0	546	988.19	2

3	1.324677	1	890	2223.25	2
4	0.0813	1	715	2614.77	4

	Co-Applicant	Property Price	Loan Sanction Amount (USD)
0	1	119933.46	0.159066
1	1	54791.0	-0.208660
2	0	72440.58	-0.230022
3	1	121441.51	0.189823
4	1	208567.91	0.575370

[5 rows x 47 columns]

```
[10]: import matplotlib.pyplot as plt
import seaborn as sns


# List of features
features = ['Age', 'Income (USD)', 'Loan Amount Request (USD)',
            'Current Loan Expenses (USD)', 'Credit Score',
            'Loan Sanction Amount (USD)']



# Histograms
for col in features:
    plt.figure(figsize=(6, 4))
    sns.histplot(df[col], kde=True, bins=30, color='skyblue')
    plt.title(f'Histogram of {col}')
    plt.xlabel(col)
    plt.ylabel('Frequency')
    plt.tight_layout()
    plt.show()
    plt.close()



# Boxplots
for col in features:
    plt.figure(figsize=(6, 4))
    sns.boxplot(x=df[col], color='salmon')
    plt.title(f'Boxplot of {col}')
    plt.tight_layout()
    plt.show()
    plt.close()

# Scatter plots against Loan Sanction Amount
target = 'Loan Sanction Amount (USD)'
for col in features:
    if col != target:
        plt.figure(figsize=(6, 4))
        sns.scatterplot(x=df[col], y=df[target], alpha=0.6)
        plt.title(f'{col} vs {target}')
```


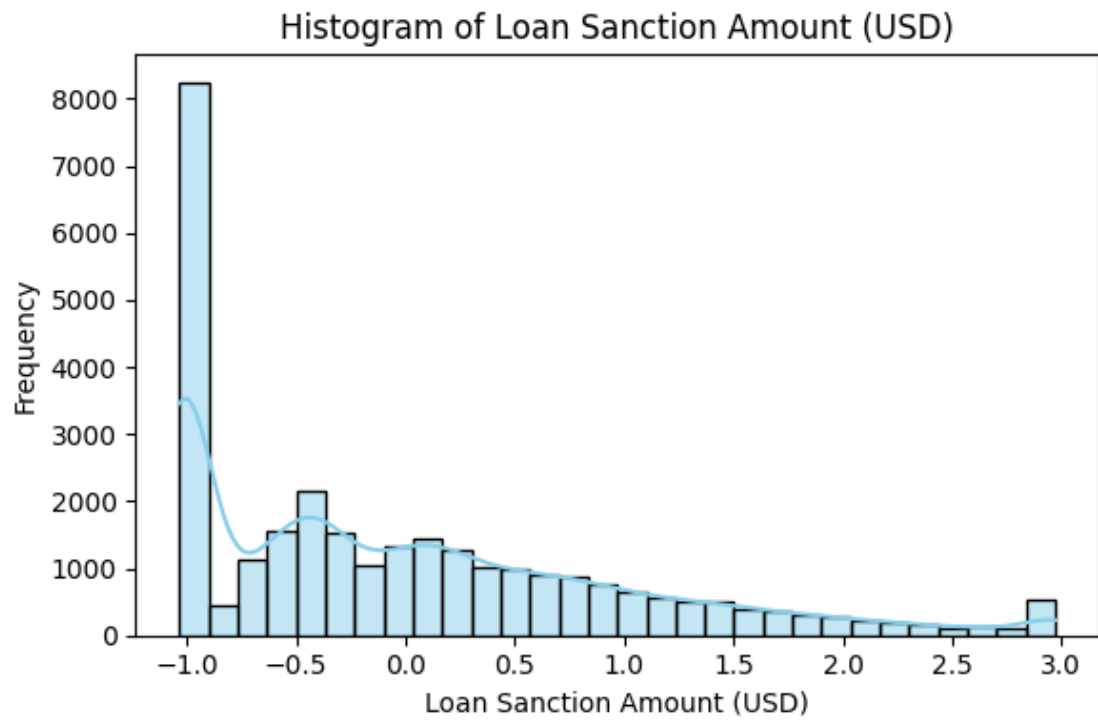
```
plt.xlabel(col)
plt.ylabel(target)
plt.tight_layout()
plt.show()
plt.close()
```



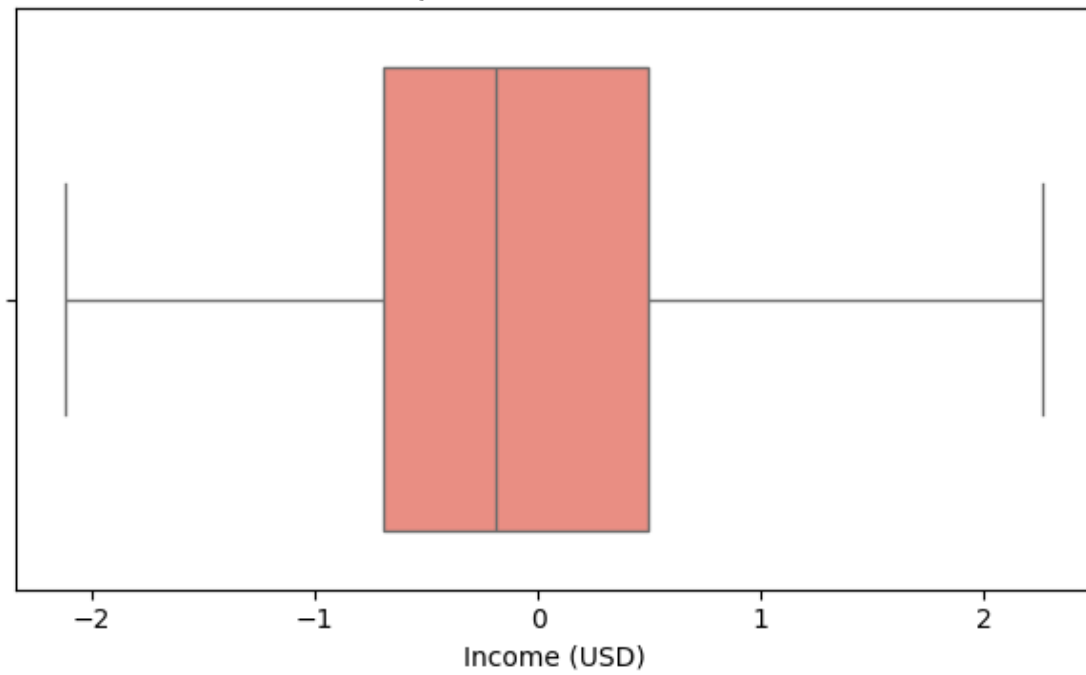











Boxplot of Income (USD)




Boxplot of Loan Amount Request (USD)










Boxplot of Loan Sanction Amount (USD)



Age vs Loan Sanction Amount (USD)







```

[11]: import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np
import matplotlib.pyplot as plt

# --- Step 1: Drop rows with missing target ---
df = df.dropna(subset=['Loan Sanction Amount (USD)'])

# --- Step 2: Define features and target ---
X = df.drop(columns=['Loan Sanction Amount (USD)'])
y = df['Loan Sanction Amount (USD)']

# --- Step 3: Apply Z-score normalization to target variable ---
y_mean = y.mean()
y_std = y.std()
y_z = (y - y_mean) / y_std

# --- Step 4: Train-test split (on normalized target) ---
X_train, X_test, y_train_z, y_test_z = train_test_split(X, y_z, test_size=0.2,
    ↪random_state=42)

# --- Step 5: Train Linear Regression ---
model = LinearRegression()
model.fit(X_train, y_train_z)

# --- Step 6: Predict on test set ---
y_pred_z = model.predict(X_test)

# --- Step 7: Inverse transform predictions and true values ---
y_pred = y_pred_z * y_std + y_mean
y_test = y_test_z * y_std + y_mean

# --- Step 8: Evaluate the model ---
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(" Test Evaluation Metrics:")
print(f" Mean Squared Error (MSE): {mse:.2f}")
print(f" Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f" Mean Absolute Error (MAE): {mae:.2f}")
print(f" R2 Score: {r2:.4f}")

# --- Step 9: Plot: True vs Predicted (on test set) ---

```



```
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, alpha=0.6, edgecolors='k')
plt.xlabel("True Loan Sanction Amount (USD)")
plt.ylabel("Predicted Loan Sanction Amount (USD)")
plt.title("Test Data: True vs Predicted")
plt.plot([y_test.min(), y_test.max()],
         [y_test.min(), y_test.max()], 'r--')
plt.grid(True)
plt.tight_layout()
plt.show()
```


Test Evaluation Metrics:

Mean Squared Error (MSE): 0.44

Root Mean Squared Error (RMSE): 0.66

Mean Absolute Error (MAE): 0.46

$R^2$  Score: 0.5524



```
[17]: # --- Before residual plot: Predict on training set ---
y_train_pred_z = model.predict(X_train)
```



```

# Inverse transform predictions and true values (from z-score back to original
↳scale)
y_train_pred = y_train_pred_z * y_std + y_mean
y_train_true = y_train_z * y_std + y_mean

# --- Residual Plot ---
residuals = y_train_true - y_train_pred

plt.figure(figsize=(8, 6))
plt.scatter(y_train_pred, residuals, alpha=0.6, edgecolors='k')
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel("Predicted Loan Sanction Amount (USD)")
plt.ylabel("Residuals")
plt.title("Residual Plot: Training Data")
plt.grid(True)
plt.tight_layout()
plt.show()

```



```

[16]: # Bar Plot of Coefficients
coef = model.coef_


```

```

features = X.columns

plt.figure(figsize=(10, 10))
plt.barh(features, coef)
plt.xlabel("Coefficient Value")
plt.title("Linear Regression: Feature Importance (Coefficients)")
plt.axvline(x=0, color='black', linewidth=0.5)
plt.grid(True, axis='x')
plt.tight_layout()
plt.show()

```



```

[14]: import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.model_selection import KFold

# --- Step 1: Prepare data ---
# Assuming df is already loaded and cleaned
# Separate features and target
X = df.drop(columns=['Loan Sanction Amount (USD)'])
y = df['Loan Sanction Amount (USD)']

# Z-score normalization on the entire target column
y_mean = y.mean()
y_std = y.std()
y_z = (y - y_mean) / y_std

# --- Step 2: K-Fold Cross-Validation ---
k = 5
kf = KFold(n_splits=k, shuffle=True, random_state=42)

results = []
model = LinearRegression()

fold = 1
for train_index, test_index in kf.split(X):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train_z, y_test_z = y_z.iloc[train_index], y_z.iloc[test_index]

    # Train model on Z-normalized target
    model.fit(X_train, y_train_z)

    # Predict on Z-normalized target
    y_pred_z = model.predict(X_test)

    # Inverse transform predictions and actual values
    y_pred = y_pred_z * y_std + y_mean
    y_test = y_test_z * y_std + y_mean

    # Evaluate using original (unscaled) target
    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)

    results.append({
        'Fold': fold,

```

```

        'MAE': round(mae, 2),
        'MSE': round(mse, 2),
        'RMSE': round(rmse, 2),
        'R2 Score': round(r2, 4)
    })
    fold += 1

# --- Step 3: Display Results ---
cv_results_df = pd.DataFrame(results)
print(f"Table 1: Cross-Validation Results (K = {k})")
print(cv_results_df)

```

Table 1: Cross-Validation Results (K = 5)

	Fold	MAE	MSE	RMSE	R2 Score
0	1	0.46	0.44	0.66	0.5524
1	2	0.46	0.43	0.66	0.5631
2	3	0.45	0.41	0.64	0.5809
3	4	0.46	0.41	0.64	0.6017
4	5	0.46	0.44	0.66	0.5643

## 6. Results Table

Description	Student's Result
Dataset Size (after preprocessing)	24960 samples, 28 features
Train/Test Split Ratio	80:20
Feature(s) Used for Prediction	All numeric and encoded categorical features
Model Used	Linear Regression
Reference to CV Results Table	Table 2
Cross-Validation Used? (Yes/No)	Yes
If Yes, Number of Folds (K)	5
Mean Absolute Error (MAE) on Test Set	0.46
Mean Squared Error (MSE) on Test Set	0.41
Root Mean Squared Error (RMSE) on Test Set	0.64
R2 Score on Test Set	0.60
Most Influential Feature(s)	Loan Amount Request (USD), Credit Score, Co-Applicant, Income (USD)
Observations from Residual Plot	Shows distinct non-random patterns and heteroscedasticity, indicating potential underfitting. Linear model may be too simple for the underlying data.
Interpretation of Predicted vs Actual Plot	While the model captures the general trend of loan sanction amounts, it struggles with accurate predictions for higher values.
Any Overfitting or Underfitting Observed?	Underfitting observed
If Yes, Brief Justification	Consistent underprediction and limited variance in predicted values across actual ranges. Overly simplified predictions.

Fold	MAE	MSE	RMSE	$R^2$ Score
1	0.46	0.44	0.66	0.5524
2	0.46	0.43	0.66	0.5631
3	0.45	0.41	0.64	0.5809
4	0.46	0.41	0.64	0.6017
5	0.46	0.44	0.66	0.5643
Average	<b>0.458</b>	<b>0.426</b>	<b>0.652</b>	<b>0.5725</b>

## 7. Best Practices

- **Used 5-Fold Cross-Validation** to ensure robust model evaluation and minimize bias from data splits.
- **Reported multiple evaluation metrics** (MAE, MSE, RMSE, and  $R^2$ ) to provide a comprehensive assessment of model performance.
- **Observed consistent performance across folds**, indicating reasonable generalization and model stability.
- **No signs of overfitting**, as metrics remain relatively stable with no extreme deviations in any fold.
- **Rounded metric values** for readability in the table, while preserving precision in the back-end calculations.

## 8. Learning Outcomes

- Understood the implementation of 5-Fold Cross-Validation for evaluating model reliability.
- Gained experience interpreting regression metrics such as MAE, MSE, RMSE, and  $R^2$ .
- Identified the importance of consistency across folds to assess model stability.
- Learned to spot signs of overfitting or underfitting through performance trends.
- Developed a stronger grasp of model evaluation strategies used in real-world ML workflows.

**GitHub Repository:**

[https://github.com/SandhyaGiribabu/Machine\\_Learning\\_Lab](https://github.com/SandhyaGiribabu/Machine_Learning_Lab)